

Mihir Kumar

+91-7428872048 | mihir.kumar@alumni.iitd.ac.in | [Linkedin](#)

EDUCATION

Indian Institute of Technology Delhi

July'16 - July'20

Bachelor's of Technology

All India Rank: *Obtained a rank 776 out of 1.5 million students across India in JEE (advanced)*

WORK EXPERIENCE

Deccan AI, Senior Software Engineer

June 2024 – Present

- Founding member of the Engineering team responsible for building core platform infrastructure and foundational backend systems. Led the data annotation product team.
- Designed and built a highly scalable, multi-stage data annotation platform supporting 30K+ daily tasks across 1000+ annotators, generating **\$1M+** monthly revenue.
- Architected a configurable, template-driven workflow engine with extensible stage types for annotation, review, LLM-based validation, proctoring, and workflow supervision, enabling new stages to be introduced without code changes. Reduced workflow onboarding TAT from weeks to hours.
- Implemented AI-based fraud detection, sampling, and proctoring mechanisms, reducing manual moderation effort by 80% while improving annotation quality and reducing fraud-induced losses.
- Structured backend systems into isolated domains including data inventory, master data, workflow orchestration, and reporting for improved maintainability and feature scalability.
- Built cloud-agnostic, self-hostable microservices enabling enterprise deployments in client-managed environments with scalable and modular architecture.
- **Google, Cisco, Snowflake, Microsoft:** Onboarded and delivered product for renowned enterprise clients.
- Built a configuration-driven form builder enabling **no-code** UI creation for annotation and review workflows, supporting rich components including text, options, markdown, images, and code blocks, significantly reducing engineering dependency for UI changes.
- Developed a template-based reporting infrastructure which generated structured reports using PostgreSQL and MongoDB data sources, reducing manual reporting effort by 99%. Enabled business and operations teams to independently generate post-annotation reports, validations, and datasets without custom engineering queries.
- Decomposed authentication and authorization into a standalone microservice consumed by 10+ internal services, supporting 100K+ user validations per day. Implemented RBAC, user blacklisting, and fraud-blocking mechanisms while eliminating reliance on third-party authentication providers.
- Built a template based communication service delivering **1M+ notifications/month** across email, WhatsApp, SMS, and in-app channels with configuration-driven provider switching for reliability and cost optimization.
- Implemented task order management and review reconciliation systems for accurate and traceable payouts.
- Ensured concurrency-safe payment processing using distributed locking (Redlock).
- Reduced tech debt and refactored critical backend systems to lower code complexity, improve maintainability, and accelerate developer productivity using AI-assisted engineering workflows and scalable software design practices.

BlinkIt, Software Engineer

Sep'20 - Oct'21

- Built backend services to automate customer retention offers and targeted notifications at scale.
- Developed and maintained Python MySQL data pipelines within JupyterHub for marketing and growth analytics.
- Automated ingestion workflows feeding Redash dashboards consumed across multiple business teams.
- Optimized batch execution and query patterns to support reliable high-volume data processing.

CitiCorp, Summer Intern, Machine Learning

May'19 - July'19

- Built and productionized a Random Forest-based risk classification model using customer transaction history.
- Delivered an end-to-end ML pipeline integrated with internal risk systems and received a Pre-Placement Offer.

SKILLS & INTERESTS

AI & GenAI: Vertex AI, LangChain, RAG, Agentic AI Frameworks, Prompt Engineering

Languages: Go, Java, Python, C/C++, Bash, React, TypeScript, JavaScript

Backend: Distributed Systems, Microservices, REST APIs, Spring Boot

Infrastructure: Git, Kafka, Prometheus, Docker, PostgreSQL, MySQL, MongoDB, SQLite, NoSQL